

Tony Yiding Tian

Philadelphia, PA | (267)249-1202 | tonytg@seas.upenn.edu | tonyxtian.com

EDUCATION

University of Pennsylvania - School of Engineering and Applied Sciences Philadelphia, PA
B.S.E. in Computer Engineering, Accelerated M.S.E. in Computer Graphics and Game Technology May 2027

- GPA: 3.8 — Relevant Courses: GPU Programming, Advanced Rendering, Interactive Computer Graphics, Computer Animation, Operating Systems, Data Structures & Algorithms, Computer Architecture

PROJECTS

-
- MatForge - Advanced Material Rendering System** | *Vulkan, Slang, C++* Nov 2025 – Dec 2025
- Architected production-quality GPU path tracer implementing four SIGGRAPH papers (2023-2024) in unified Vulkan ray tracing pipeline, contributing 2,400+ lines of C++ and Slang code for sampling and geometry systems
 - Implemented Quad-Optimized Low-Discrepancy Sequences (QOLDS) from SIGGRAPH 2024, engineering base-3 Sobol' generator with Owen scrambling across 47 dimensions, achieving +2.57 dB PSNR and 44.7% MSE reduction vs PCG sampling at 512 SPP with < 1% performance overhead
 - Developed RMIP (Rectangular MinMax Image Pyramid) intersection shader, enabling tessellation-free displacement mapping through hierarchical texture-space ray traversal with custom Vulkan intersection shaders
 - Integrated techniques into complete Monte Carlo path tracing pipeline supporting glTF 2.0 scenes, HDR environment maps, and KHR_materials_displacement extension for physically-based rendering
 - GitHub Repo: github.com/MatForge/MatForge
- Vulkan Grass Renderer** | *Vulkan, GLSL, C++* Oct 2025 – Nov 2025
- Implemented real-time grass rendering system in Vulkan, simulating up to 1 million grass blades with physics-based animation at interactive frame rates using compute and tessellation pipelines
 - Developed GPU compute shader for physics simulation applying gravity, Hooke's law recovery forces, and wind with spatial turbulence to quadratic Bezier curve grass blades, processing 1M+ blade state updates per frame
 - Engineered three-tiered GPU culling system (orientation, view-frustum, distance-based probabilistic) achieving 4.31× performance improvement at 1M blades through efficient atomic operations and indirect draw calls
 - Built hardware tessellation pipeline with dynamic LOD based on camera distance, generating smooth blade geometry with exponential falloff from 10 to 2 subdivisions for optimal geometry density
 - Created comprehensive ImGui control panel with real-time FPS metrics (average, 1% lows), frame time graphs, and automated performance testing system generating CSV benchmarks across 40 configurations
 - GitHub Repo: github.com/tonygrt/Vulkan-Grass-Renderer
- CUDA Path Tracer** | *CUDA, GLSL, C++* Sep 2025 - Oct 2025
- Monte Carlo path tracer capable of rendering complex 3D scenes with custom 3D models and environment maps
 - Implemented shading BSDF kernel supporting global illumination, multiple importance sampling, anti-aliasing, sub-surface scattering, capable of rendering various PBR material types with albedo and texture maps
 - Integrated third-party libraries of tinyGLTF to support glTF 2.0 mesh loading and Nvidia OptiX for denoising
 - Utilized various techniques to boost performance: material sorting (+5%), Russian Roulette (+6% - 24%), stream compaction (+24% - 67%), and Bounding Volume Hierachy (3× - 160× framerate in complex scenes)
 - Project Repo and Demo: github.com/tonygrt/CUDA-Path-Tracer. A previous standalone performance focused stream compaction project with detailed analysis in Nsight: github.com/tonygrt/Project2-Stream-Compaction
- WebGPU Renderer** | *WebGPU, TypeScript, WGSL* Sep 2025 – Oct 2025
- Implemented three advanced rendering techniques for real-time lighting of 5000+ dynamic point lights: Naive Forward, Forward+, and Clustered Deferred rendering using WebGPU compute and graphics pipelines
 - Engineered screen-space light clustering system using compute shaders, subdividing view frustum into 16×9×24 grid with exponential depth slicing and sphere-AABB intersection testing for efficient light culling
 - Developed G-buffer architecture with 3 render targets (position, normal, albedo) enabling two-pass deferred rendering that decouples geometry complexity from lighting calculations
 - Achieved 53x performance improvement over naive rendering (497ms to 9.3ms at 5000 lights) and 3.5x speedup vs Forward+ through overdraw elimination and optimized memory access patterns
 - Built automated performance testing system collecting statistical frame time data across 30 configurations (3 renderers × 10 light counts), generating CSV analysis for rigorous benchmarking
 - Live demo deployed at webgpu.tonyxtian.com, rendering Sponza scene with real-time performance metrics and interactive controls. GitHub repo: github.com/tonygrt/Project4-WebGPU-Forward-Plus-and-Clustered-Deferred

TECHNICAL SKILLS

Programming: C++, CUDA, Python, WGSL, GLSL, Parallel algorithms, Memory management, Rendering pipeline
Graphics/Rendering: Nsight Profiling, Path Tracing, Deferred Rendering, Rasterization, Animation systems, PBR
Tools/APIs: Nvidia Nsight, WebGPU, Vulkan, Visual Studio, Qt, OpenGL, Git, CMake, MakeFile, Clang, GDB